

1 Minimum Spanning Trees

Suppose we're given a connected, undirected, weighted graph $G = (V, E)$, and a function $w : e \rightarrow \mathbb{R}$ that returns the weight of an edge $e \in E$. As we discussed last time, a *minimum spanning tree* (MST) is a spanning tree T that minimizes the function

$$w(T) = \sum_{e \in T} w(e) .$$

To simplify our analysis, let us assume that $\forall_{e, e' \in E} w(e) \neq w(e')$, that is, the edge weights are distinct. This assumption is convenient because it guarantees that the minimum spanning tree of the graph will be unique.¹ (Can you prove this is true?)

If there are ties among some of the edge weights, then there may be more than one MST. As an extreme example, consider the case where all the edges have unit weight, i.e., $\forall_{e \in E} w(e) = 1$. In this case, not only is there no unique MST, but every spanning tree (including every SSSP) is a MST and the weight of every spanning tree is exactly $w(T) = V - 1$.

What good are MSTs? Generally, they are only useful if you need a subgraph T that connects all pairs within the graph G and you need that subgraph to have minimum weight. For instance, suppose we have a set of points embedded in space, e.g., cities in a country, homes in a city, sea ports worldwide, etc., and we want to build a distribution “network” that allows fast transportation of some good, e.g., freight cargo, electricity over wires, plush toys to market, etc., between the different points. If cost is no object, then the best solution is to connect every pair of points with an edge, i.e., a fully connected graph G on the n points. This is also the solution with the largest cost, and our shareholders or clients might prefer to exchange slightly less efficient routing for a lower cost solution. A MST provides a minimum-cost solution to this problem that still allows routing between any pair of points. MSTs can also be used as approximate solutions to genuinely hard problems like the Traveling Salesman Problem (TSP), which is NP-Hard.²

As an example, Figure 1 shows the MST and SSSP from Lecture 7.

¹More generally, this implies that the function $w(T)$ has a single global minimum over the space of all trees T on G . It does not, however, say anything about the number of alternative trees that are very close in weight to the global minimum, or how similar these alternative trees are to the global optimum. Normally, the number and dissimilarity of these approximately-good solutions is irrelevant, but in applications where our algorithm is not guaranteed to always find the best solution, as with NP-Hard problems, or when it's not clear that the minimum-weight solution is the best solution for other reasons, alternatives can matter a lot.

²As in the case of the Christofides algorithm for approximating TSP.




Figure 1: Recall from Lecture 7 that a single-source shortest path (SSSP) tree and a minimum spanning tree are not necessarily the same. Here's the example I gave to illustrate that point. The MST is on the left, and the SSSP is on the right.

1.1 A generic MST algorithm

As with the search-tree problem (Lecture 6) and the single-source shortest path problem (Lecture 7), different MST algorithms are in fact simply instances of a generic spanning tree algorithm.

This generic algorithm maintains an acyclic subgraph F on the input graph G , called an *intermediate spanning forest*. The forest F is a subgraph of the MST of G , and every component of F is a MST on that component's vertices. Initially, F has no edges and so is just a collection of V disconnected singleton components; when the algorithm halts, F is a single V -node tree, which is a MST of G . At each step of the algorithm, it adds a single edge to F , but to guarantee the halting condition, we must be careful about which edge we choose to add.

At each intermediate step of the algorithm, F induces two special types of edges, as well as a third less-special type. Call an edge *useless* if it is not an edge in F but both of its endpoints are in the same component of F . Adding a useless edge to F would induce a cycle, which is why we call them useless. In contrast, for each component of F , an edge is called *safe* if it is the minimum weight edge among all edges with exactly one endpoint in that component. Because we assumed that edge weights are distinct, a given component can only have one safe edge.

Adding a safe edge to F increases the size of the spanning tree and maintains its minimum weight, which is why we call them safe. (If two components in F have the same safe edge, what happens next?) Edges that are neither safe nor useless are called *undecided*.

From these observations, we can write down pseudo-code for our generic algorithm:

```

Generic-MST( $G, w$ ) {
     $F = \{\}$                                 // initially  $F$  contains no edges
    while ( $F$  not a spanning tree) {        // are we done yet?
        find edge  $(u, v)$  that is safe for  $F$  // make a safe choice
        add  $(u, v)$  to  $F$                     //
    }
    return  $F$                                 // we're done
}

```

That is, at each step of the algorithm, we identify a safe edge and add it to the intermediate spanning forest. When we add an edge to F , some undecided edges become safe while others become useless. We stop when F is a spanning tree. (Or, we stop when there are no more safe edges. Do you see why these are equivalent?)

Lemma 1: A minimum spanning tree contains all safe edges and no useless edges.

Proof: Let T' be a minimum spanning tree. Suppose F has a “bad” component in it, whose safe edge $e = (u, v)$ does not end up in T' . Since T' is connected, it contains a unique path from u to v and therefore there is at least one edge e' on this path that has exactly one endpoint in the bad component. If we replace e' in T' with e , we get a new spanning tree T . By assumption, e is the bad component’s safe edge, which implies that $w(e') > w(e)$ and that the new spanning tree T has smaller weight than T' . However, this contradicts our assumption that T' is a MST; therefore, T' must, in fact, already contain every safe edge and e cannot be a safe edge. Further, if we added any useless edge to F , we would introduce a cycle. \square




Figure 2: Every safe edge is in the MST. The “bad” component in our proof of Lemma 1 is the left component (highlighted in blue). The path from u to v on T' is given in bold and traverses e' .

1.1.1 Borůvka's algorithm

One of the earliest solutions to the MST problem is due to Borůvka (1926).³ Unlike the more commonly known versions, this algorithm does not add safe edges one at a time. In Borůvka's version of **Generic-MST**, F is again a forest. The critical choice is that *we add all current safe edges simultaneously and then recurse until there are no safe edges left*.

Each recursion is a “phase” and during each phase the algorithm first elects an arbitrary “leader” for each component (it does not matter which vertex is a leader, just that each component has one) and then identifies the set of safe edges. There are several ways to choose the leaders; a simple one is to use the **Search-Forest** algorithm (the generalization of **Search-Tree** that builds a forest of search-trees on G), e.g., DFS.

Here is pseudo-code for Borůvka's algorithm (assume that $w(e) = +\infty$ if $e \notin E$):

```

Boruvka(G,w) {
  F = {V,{}}
  while F has more than one component
    choose component leaders      % how long does this take?
    Find-Safe-Edges(G,w)         % find all the currently safe edges
    for each leader v' {         % for each component
      add safe(v') to F          % add its safe edge to F
    }
}

Find-Safe-Edges(G,w) {
  for each leader v' {           % for each component,
    safe(v') = NULL              % no safe edge
  }
  for each edge (u,v) in E {
    u' = leader(u)               % look up component name of u
    v' = leader(v)               % look up component name of v
    if u' != v' {                % are u, v in same component
      if w(u,v) < w(safe(u')) { safe(u') = (u,v) } % update estimated safe edge
      if w(u,v) < w(safe(v')) { safe(v') = (u,v) } %
    }
  }
}

```

Note that **Find-Safe-Edges** is essentially a brute-force search to identify which edges are safe for each component. Because it examines every edge, calling it takes $O(E)$ time, if the edges are stored in an adjacency list format. (How long if we use an adjacency matrix?) Because we add one safe

³Depending on the literature, this algorithm is also called Choquet's algorithm, Florek-Lukaziewicz-Perkal-Stienhaus-Zubrzycki's algorithm or Sollin's algorithm, all of whom reinvented it after Borůvka's original 1926 paper.

edge to F for each component, and every safe edge joins two components, each pass through the main loop reduces the number of components by at least half. (Do you see how it could be more than half?) Since there are initially V components, this yields $O(\log V)$ passes through the main loop, each of which takes $O(E)$ time. Thus, the running time of Borůvka's algorithm is $O(E \log V)$.

One nice thing about Borůvka's algorithm is that it is a naturally parallel algorithm, in the sense that the identification of safe edges can be done in parallel for each of the components.

Figure 3 shows an example of running Borůvka's algorithm on a small graph.




Figure 3: Borůvka's algorithm run on the example graph. The thick (red) edges are in F and the (blue) arrows indicate each component's safe edges. Useless edges are shown as dashed lines. (Figure from Jeff Erickson's lecture notes.)

1.1.2 Jarník's (Prim's) algorithm

The second oldest MST algorithm, and one that typically appears in textbooks, is originally due to Jarník (1929), but usually bears Prim's name, who rediscovered it in 1957.⁴ In Prim's algorithm, **Generic-MST** maintains F as a forest, composed of one tree, which we call T , containing S vertices and $V - S$ singleton components. The critical choice is then to *find T 's safe edge and add it to T* . In this way, we grow the MST out from some initial vertex until it spans the graph.

Initially, T contains only this arbitrarily selected vertex. The algorithm then grows T outward by repeatedly finding and adding T 's safe edge. When the algorithm halts, T spans the graph. To implement Prim's algorithm, we need a fast way to find the current safe edge. This can be done using a min-heap of the edges adjacent to T , using the edge weights as the keys. When we pop the minimum-weight edge from the top of the heap, we need to check whether it is a useless edge

⁴It was also discovered by Kruskal, Loberman and Weinberger, and by Dijkstra.

(both end-points in T) or not. If not, then it is a safe edge and we then both add it to T and add the edges attached to the new member of T to the heap.

Note that because Prim's algorithm grows the MST, it is also a version of the **Generic-SSSP** algorithm from Lecture 7, albeit one that uses the min-heap as the "set" ADT (this makes it equivalent to Dijkstra's algorithm). As a result, the running time is $O(E \log E) = O(E \log V)$.

Figure 4 shows running Prim's algorithm on the example graph.




Figure 4: Prim's algorithm run on the example graph, starting with the bottom-most vertex. Again, thick (red) edges are in F , (blue) arrows indicate the safe edge, and dashed lines show the useless edges. (Figure from Jeff Erickson's lecture notes.)

But note that Prim's algorithm only visits each vertex once. This means that we can speed up the running time of this algorithm by being smarter about how we find safe edges. Instead of maintaining a min-heap on the edges attached to T , we can maintain a heap of the vertices themselves, where the key of a vertex v is the length of the minimum-weight edge between v and T (or ∞ if we haven't seen an edge yet). Now, when we add a new edge to T , we might need to decrease the keys of its neighbors.

In this version of Prim's algorithm, the running time is dominated by the heap operations, and in particular, the cost of inserting all V keys into the initial heap, the cost of getting the minimum V times, and the cost of modifying the keys. The first and second of these take $O(V)$ times, and at

most, each edge induces a key modification, so there are $O(E)$ of those. If we use a standard binary heap, each operation costs $O(\log V)$, so the overall running time would be $O((V + E) \log V) = O(E \log V)$. We can do slightly better if we use the Fibonacci heap, which yields a running time of $O(E + V \log V)$.

1.1.3 Kruskal's algorithm

Another popular (that is, it appears in many textbooks) MST algorithm is named for Kruskal (1956). In this version of **Generic-MST**, F is maintained as a forest. The critical choice is that, at each step of the algorithm, we *add to F the smallest weight edge that is also safe*.

That is, we sort the edges by their weights and then (repeatedly) scan them in order; we add the first safe edge we find to F and then start over. Recall that a safe edge is one that has its endpoints in different components. To determine if an edge (u, v) is safe, we use the same approach as Borůvka, i.e., each component is assigned a “leader” and each vertex keeps track of which leader it is following. Unlike Borůvka’s algorithm, however, we do not recompute leaders each time we add an edge; instead, we simply let the old leader follow the new leader (formally, this is equivalent to using the Union-Find algorithm to keep track of the leader hierarchy; we’ll see more about Union-Find in a few weeks). Here’s pseudo-code:

```
Kruskal(G,w) {
    sort E by weight
    F = {}
    for each vertex v in V {
        Make-Set(v)                // initialize Union-Find
    }
    for i = 1 to E {
        (u,v) = ith lightest edge in E
        if Find(u) != Find(v) {    // different leaders?
            Union(u,v)             // merge components
            add (u,v) to F
        }
    }
}
```

This algorithm performs $O(E)$ Find operations, one for each end-point of each edge in the graph, and $O(V)$ Union operations, one for each edge in the spanning tree. We haven’t talked about how long these operations take: they take $O(\alpha(E, V))$ time, where $\alpha(.,.)$ is the inverse-Ackerman function. The main thing you need to know about this function is that it grows extremely slowly⁵ and can, for most purposes be treated as a constant function. Thus, the overall running time is

⁵Slower, in fact, than any function you can express in standard algebraic notation. If this sounds crazy to you, and it should, you should take a look at Scott Aaronson’s short and highly entertaining essay “Who Can Name The Bigger Number?” which you can find here: <http://www.scottaaronson.com/writings/bignumbers.html>

dominated not by the **Union** and **Find** operations, but by the time to sort the edges by their weight, which takes $O(E \log E) = O(E \log V)$ time.




Figure 5: Kruskal's algorithm run on the example graph. Again, thick (red) edges are in F and dashed lines are useless edges. (Figure from Jeff Erickson's lecture notes.)

2 For Next Time

1. Read Chapter 26 (Max Flow Algorithms)